

Synchronous single initiator spanning tree algorithm using flooding

Siddharth Bhat, Anurag Chaturvedi, Hitesh Kaushik

March 13, 2020

Introduction

- ▶ We use BFS to compute a spanning tree of a graph.


Introduction

- ▶ We use BFS to compute a spanning tree of a graph.
- ▶ We can distribute the sequential algorithm assuming *synchronous communication*.


Introduction

- ▶ We use BFS to compute a spanning tree of a graph.
- ▶ We can distribute the sequential algorithm assuming *synchronous communication*.
- ▶ **Refresher:** Algorithm proceeds in rounds. All messages sent at round i are received at round i .


Introduction

- ▶ We use BFS to compute a spanning tree of a graph.
- ▶ We can distribute the sequential algorithm assuming *synchronous communication*.
- ▶ **Refresher:** Algorithm proceeds in rounds. All messages sent at round i are received at round i .
- ▶ **Flood:** everyone sends a message to their neighbours. Think flood fill. 


Example




Example




Example




Example




Example




Example




Example




Why Synchrony?




Why Synchrony?




Why Synchrony?




Why Synchrony?




Why Synchrony?



Why Synchrony?



Why Synchrony?



Synchronous BFS (Pseudocode)

- ▶ Assume root begins computation.
- ▶ Algorithm is synchronous.

Synchronous BFS (Pseudocode)

- ▶ Assume root begins computation.
- ▶ Algorithm is synchronous.

```
def bfs_spanning_tree(self):
```

Synchronous BFS (Pseudocode)

- ▶ Assume root begins computation.
- ▶ Algorithm is synchronous.

```
def bfs_spanning_tree(self):  
    if self.id == ROOT_ID:  
        self.visited = True; self.depth = 0;  
        for n in self.neighbours: n.send(self.id)
```

Synchronous BFS (Pseudocode)

- ▶ Assume root begins computation.
- ▶ Algorithm is synchronous.

```
def bfs_spanning_tree(self):  
    if self.id == ROOT_ID:  
        self.visited = True; self.depth = 0;  
        for n in self.neighbours: n.send(self.id)  
  
    for round in range(1, DIAMETER+1):  
        if not self.visited: # if visited, skip
```

Synchronous BFS (Pseudocode)

- ▶ Assume root begins computation.
- ▶ Algorithm is synchronous.

```
def bfs_spanning_tree(self):  
    if self.id == ROOT_ID:  
        self.visited = True; self.depth = 0;  
        for n in self.neighbours: n.send(self.id)  
  
    for round in range(1, DIAMETER+1):  
        if not self.visited: # if visited, skip  
  
            if self.queries: # if we have a query  
                # randomly choose from queries  
                parent = random.choice(self.query)  
                self.visited = True  
                self.depth = round
```


Synchronous BFS (Pseudocode)

- ▶ Assume root begins computation.
- ▶ Algorithm is synchronous.

```
def bfs_spanning_tree(self):  
    if self.id == ROOT_ID:  
        self.visited = True; self.depth = 0;  
        for n in self.neighbours: n.send(self.id)  
  
    for round in range(1, DIAMETER+1):  
        if not self.visited: # if visited, skip  
  
            if self.queries: # if we have a query  
                # randomly choose from queries  
                parent = random.choice(self.query)  
                self.visited = True  
                self.depth = round  
  
                # synchronous  
                for n in self.neighbours: n.send(self.id)  
            self.queries = [];
```

Synchronous BFS (Ending earlier if visited)

```
def bfs_spanning_tree(self):  
    if self.id == ROOT_ID:  
        self.depth = 0;  
        for n in self.neighbours: n.send(self.id)  
  
        return # early-exit for root node  
  
    for round in range(1, DIAMETER+1):  
        if self.queries: # if we have a query  
            # randomly choose from queries  
            parent = random.choice(self.query)  
            self.visited = True  
            self.depth = round  
  
            # synchronous  
            for n in self.neighbours: n.send(self.id)  
  
            return # early-exit for child
```

Synchronous BFS (Learning children)

- ▶ Assume root begins computation.
- ▶ Algorithm is synchronous.

```
def bfs_spanning_tree(self):  
    if self.id == ROOT_ID:  
        self.visited = True; self.depth = 0;  
        for n in self.neighbours: n.send(self.id)  
    for round in range(1, DIAMETER+1):  
        if self.visited: # if visited, wait for children  
            for q in self.queries: self.children.append(q)  
        else: # if not visited, run code  
            if self.queries: # if we have a query  
                # randomly choose from queries  
                parent = random.choice(self.query)  
                self.visited = True  
                self.depth = round  
  
                # synchronous  
                for n in self.neighbours: n.send(self.id)  
                parent.send(self.id) # send to parent  
            self.queries = [];
```

Complexities

Complexities

- ▶ Local space for a : $|\{v : (a, v) \in E\}|$ (#of incident edges)

Complexities

- ▶ Local space for a : $|\{v : (a, v) \in E\}|$ (#of incident edges)
- ▶ $|Diameter|$ rounds.

Complexities

- ▶ Local space for a : $|\{v : (a, v) \in E\}|$ (#of incident edges)
- ▶ $|Diameter|$ rounds.
- ▶ 1 or 2 messages / edge. Message complexity $\leq 2|E|$.

Thank you!

Asynchronous Bounded Delay Network

- ▶ All processes have physical clocks: **need not be synchronized.**
- ▶ Message delivery time is bounded by constant $\mu \in \mathbb{R}$.

ABD Synchronizers: Bounded Delay \rightarrow Synchronized

- ▶ All processes have physical clocks: **need not be synchronized.**
- ▶ Message delivery time is bounded by constant $\mu \in \mathbb{R}$.

Key idea

Chunk "real time" into units of μ . Each μ block of time behaves like a logical synchronized tick!

Formal definition of synchronous processes p

► $p \equiv (S, I, M\mathcal{G}, \vdash)$

Formal definition of synchronous processes p

- ▶ $p \equiv (S, I, M\mathcal{G}, \vdash)$
- ▶ $S \equiv$ set of states.

Formal definition of synchronous processes p

- ▶ $p \equiv (S, I, M\mathcal{G}, \vdash)$
- ▶ $S \equiv$ set of states.
- ▶ $I \subseteq S \equiv$ set of initial states.

Formal definition of synchronous processes p

- ▶ $p \equiv (S, I, MG, \vdash)$
- ▶ $S \equiv$ set of states.
- ▶ $I \subseteq S \equiv$ set of initial states.
- ▶ $MG : S \rightarrow 2^M \equiv$ message generating **function**.

Formal definition of synchronous processes p

- ▶ $p \equiv (S, I, MG, \vdash)$
- ▶ $S \equiv$ set of states.
- ▶ $I \subseteq S \equiv$ set of initial states.
- ▶ $MG : S \rightarrow 2^M \equiv$ message generating **function**.
- ▶ $\vdash \subseteq S \times 2^M \subset S \equiv$ transition **relation**. We write $(s, M) \vdash s'$ for $(s, M, s') \in \vdash$.

Formal definition of synchronous processes p

- ▶ $p \equiv (S, I, MG, \vdash)$
- ▶ $S \equiv$ set of states.
- ▶ $I \subseteq S \equiv$ set of initial states.
- ▶ $MG : S \rightarrow 2^M \equiv$ message generating **function**.
- ▶ $\vdash \subseteq S \times 2^M \subset S \equiv$ transition **relation**. We write $(s, M) \vdash s'$ for $(s, M, s') \in \vdash$.
- ▶ Process starts at state $s_0 \in I$.

Formal definition of synchronous processes p

- ▶ $p \equiv (S, I, MG, \vdash)$
- ▶ $S \equiv$ set of states.
- ▶ $I \subseteq S \equiv$ set of initial states.
- ▶ $MG : S \rightarrow 2^M \equiv$ message generating **function**.
- ▶ $\vdash \subseteq S \times 2^M \subset S \equiv$ transition **relation**. We write $(s, M) \vdash s'$ for $(s, M, s') \in \vdash$.
- ▶ Process starts at state $s_0 \in I$.
- ▶ in state s , process sends a collection of messages $MG(s)$.

Formal definition of synchronous processes p

- ▶ $p \equiv (S, I, MG, \vdash)$
- ▶ $S \equiv$ set of states.
- ▶ $I \subseteq S \equiv$ set of initial states.
- ▶ $MG : S \rightarrow 2^M \equiv$ message generating **function**.
- ▶ $\vdash \subseteq S \times 2^M \subset S \equiv$ transition **relation**. We write $(s, M) \vdash s'$ for $(s, M, s') \in \vdash$.
- ▶ Process starts at state $s_0 \in I$.
- ▶ in state s , process sends a collection of messages $MG(s)$.
- ▶ **Before** next pulse, it receives *all* messages M sent to it in this pulse.

Formal definition of synchronous processes p

- ▶ $p \equiv (S, I, MG, \vdash)$
- ▶ $S \equiv$ set of states.
- ▶ $I \subseteq S \equiv$ set of initial states.
- ▶ $MG : S \rightarrow 2^M \equiv$ message generating **function**.
- ▶ $\vdash \subseteq S \times 2^M \subset S \equiv$ transition **relation**. We write $(s, M) \vdash s'$ for $(s, M, s') \in \vdash$.
- ▶ Process starts at state $s_0 \in I$.
- ▶ in state s , process sends a collection of messages $MG(s)$.
- ▶ **Before** next pulse, it receives *all* messages M sent to it in this pulse.
- ▶ It then enters a state s' such that $(s, M) \vdash s'$.

Formal definition of synchronous processes p

- ▶ $p \equiv (S, I, MG, \vdash)$
- ▶ $S \equiv$ set of states.
- ▶ $I \subseteq S \equiv$ set of initial states.
- ▶ $MG : S \rightarrow 2^M \equiv$ message generating **function**.
- ▶ $\vdash \subseteq S \times 2^M \subset S \equiv$ transition **relation**. We write $(s, M) \vdash s'$ for $(s, M, s') \in \vdash$.
- ▶ Process starts at state $s_0 \in I$.
- ▶ in state s , process sends a collection of messages $MG(s)$.
- ▶ **Before** next pulse, it receives *all* messages M sent to it in this pulse.
- ▶ It then enters a state s' such that $(s, M) \vdash s'$.
- ▶ Note that MG is **deterministic**. All the nondeterminism is in the relation $(s, M) \vdash s'$.

Formal definition of synchronous system \mathbb{P}

- ▶ $\mathbb{P} \equiv (p_1, p_2, \dots, p_n)$ where each $p_i \equiv (S_i, l_i, M\mathcal{G}_i, \vdash_i)$ is a synchronous process, is a synchronous system.

Formal definition of synchronous system \mathbb{P}

- ▶ $\mathbb{P} \equiv (p_1, p_2, \dots, p_n)$ where each $p_i \equiv (S_i, l_i, M\mathcal{G}_i, \vdash_i)$ is a synchronous process, is a synchronous system.
- ▶ Transition system $Tr \equiv (S, \mathcal{I}, \rightarrow)$.

Formal definition of synchronous system \mathbb{P}

- ▶ $\mathbb{P} \equiv (p_1, p_2, \dots, p_n)$ where each $p_i \equiv (S_i, l_i, M\mathcal{G}_i, \vdash_i)$ is a synchronous process, is a synchronous system.
- ▶ Transition system $Tr \equiv (S, \mathcal{I}, \rightarrow)$.
- ▶ $S \equiv \prod_i S_i$ is the state space of the transition system.

Formal definition of synchronous system \mathbb{P}

- ▶ $\mathbb{P} \equiv (p_1, p_2, \dots, p_n)$ where each $p_i \equiv (S_i, l_i, M\mathcal{G}_i, \vdash_i)$ is a synchronous process, is a synchronous system.
- ▶ Transition system $Tr \equiv (S, \mathcal{I}, \rightarrow)$.
- ▶ $S \equiv \prod_i S_i$ is the state space of the transition system.
- ▶ $\mathcal{I} \equiv \prod_i l_i$ is the set of initial states.

Formal definition of synchronous system \mathbb{P}

- ▶ $\mathbb{P} \equiv (p_1, p_2, \dots, p_n)$ where each $p_i \equiv (S_i, l_i, M\mathcal{G}_i, \vdash_i)$ is a synchronous process, is a synchronous system.
- ▶ Transition system $Tr \equiv (S, \mathcal{I}, \rightarrow)$.
- ▶ $S \equiv \prod_i S_i$ is the state space of the transition system.
- ▶ $\mathcal{I} \equiv \prod_i l_i$ is the set of initial states.
- ▶ $\rightarrow \subseteq S \times S$ is the transition relation.

$$(s_1, s_2, \dots, s_n) \rightarrow (s'_1, s'_2, \dots, s'_n) \\ \text{iff } \forall p_i \in \mathbb{P}; (s_i, \{\text{messages to } p_i \text{ from } p_{-i}\}) \vdash s'_i$$

